

# Database Systems

---

## *Course Project Instruction*

Database group

刘思彤 张晓航

# The Task is

To Implement a DBMS Prototype.

# Policy

- 2-3 persons form a team.
- 60% of your final score.

# DBMS: Problems

## ➤DBMS涉及到的问题

- ✓ 数据库管理系统是为了管理大量、复杂的数据。对数据的管理既涉及到
  - 存：数据存储结构的定义：存得好！
  - 取：数据操作机制的提供：取得快！
- ✓ 如果数据被多用户共享，那么DBMS还必须设法避免可能产生的异常结果，即并发控制；
- ✓ 如果系统发生故障，那么DBMS必须保证将数据恢复到故障发生前的状态，即故障恢复；
- ✓ DBMS还必须保证所存储数据的安全性，不被非法访问和操作，即访问控制……

# What We Care

- Correctness
- Response Time
  - Storage
  - Index
  - Caching Strategy
  - Optimizing
  - Query Processing

名称


|   |
|---|
|  第01章 : 概论与引言.ppt        |
|  第02章 : E-R模型.ppt        |
|  第03章 : 关系模型.ppt         |
|  第04章 : SQL.ppt          |
|  第05章 : 完整性约束.ppt        |
|  第06章 : 数据库的物理设计.ppt     |
|  第07章 : 存储结构和文件结构.ppt    |
|  第08章 : 索引和散列.ppt        |
|  第09章 : 查询处理.ppt         |
|  第10章 : 事务.ppt           |
|  第11章 : 并发控制.ppt        |
|  第12章 : 数据库系统的体系结构.ppt |
|  第13章 : 数据仓库.ppt       |

# Grading Criteria

|                    |   |          |
|--------------------|---|----------|
| Accomplishment     | At least one correct run  | 10       |
| Overall Evaluation | Correctness & Design &<br>Code Quality & Contrib.   | 10       |
| Performance        | $S_j = \text{sum}((T_{i,\text{best}} / T_{i,j})^{0.2})$<br>Full * ( $S_j / S_{\text{best}}$ ) | 30       |
| Documentation      | Content & Feature   | 10       |
| Presentation       | For some teams only   | $\leq 5$ |

# Example

|        | Workload 0 | Workload 1 |
|--------|------------|------------|
| Team 0 | 5          | 100        |
| Team 1 | 10         | 1000       |
| Team 2 | 1          | Fail       |



$$S_0 = (1 / 5)^{0.2} + (100 / 100)^{0.2} = 1.725$$

$$S_1 = (1 / 10)^{0.2} + (100 / 1000)^{0.2} = 1.262$$

$$S_2 = (1 / 1)^{0.2} + (100 / \text{INF})^{0.2} = 1$$

$$\text{Score}_0 = 30 * 1.725 / 1.725 = 30$$

$$\text{Score}_1 = 30 * 1.262 / 1.725 = 22$$

$$\text{Score}_2 = 30 * 1 / 1.725 = 17$$

$$(1 / 5)^{0.2} = 0.725$$

$$(1 / 10)^{0.2} = 0.631$$

$$(1 / 50)^{0.2} = 0.457$$

$$(1 / 100)^{0.2} = 0.398$$

$$(1 / 500)^{0.2} = 0.289$$

$$(1 / 1000)^{0.2} = 0.251$$

$$(1 / 5000)^{0.2} = 0.182$$

# The Environment is

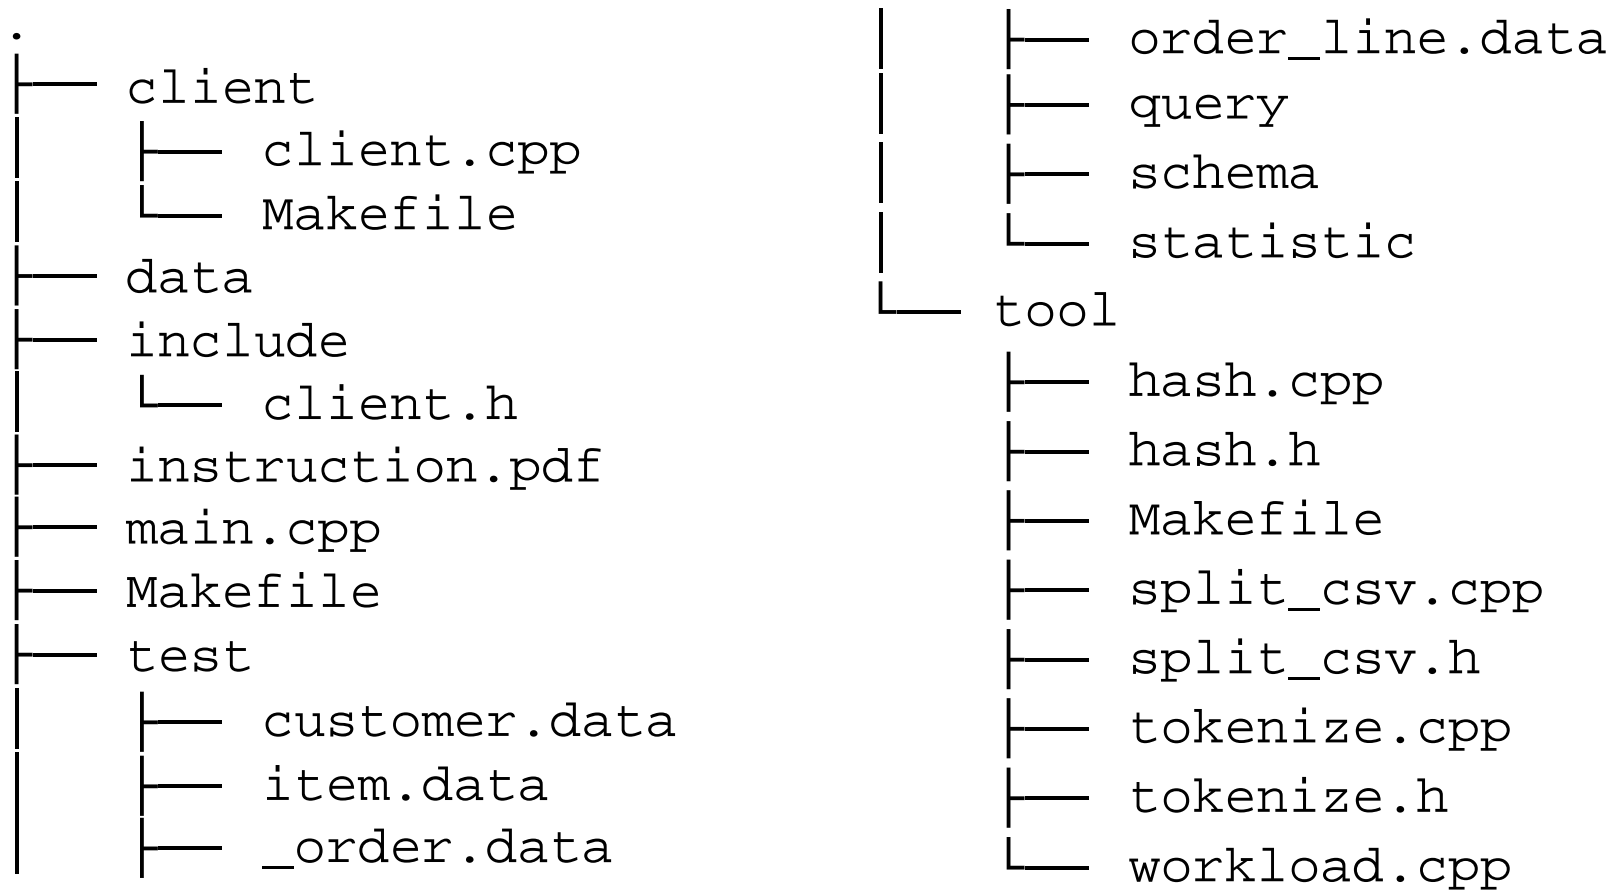
- Ubuntu 08.04 LTS 32bit/ 2CPU\*4 cores
- g++ 4.4.3
- Intel(R) Xeon(R) CPU E5420@2.50GHz
- 16GB RAM, 512GB\*2 Disk



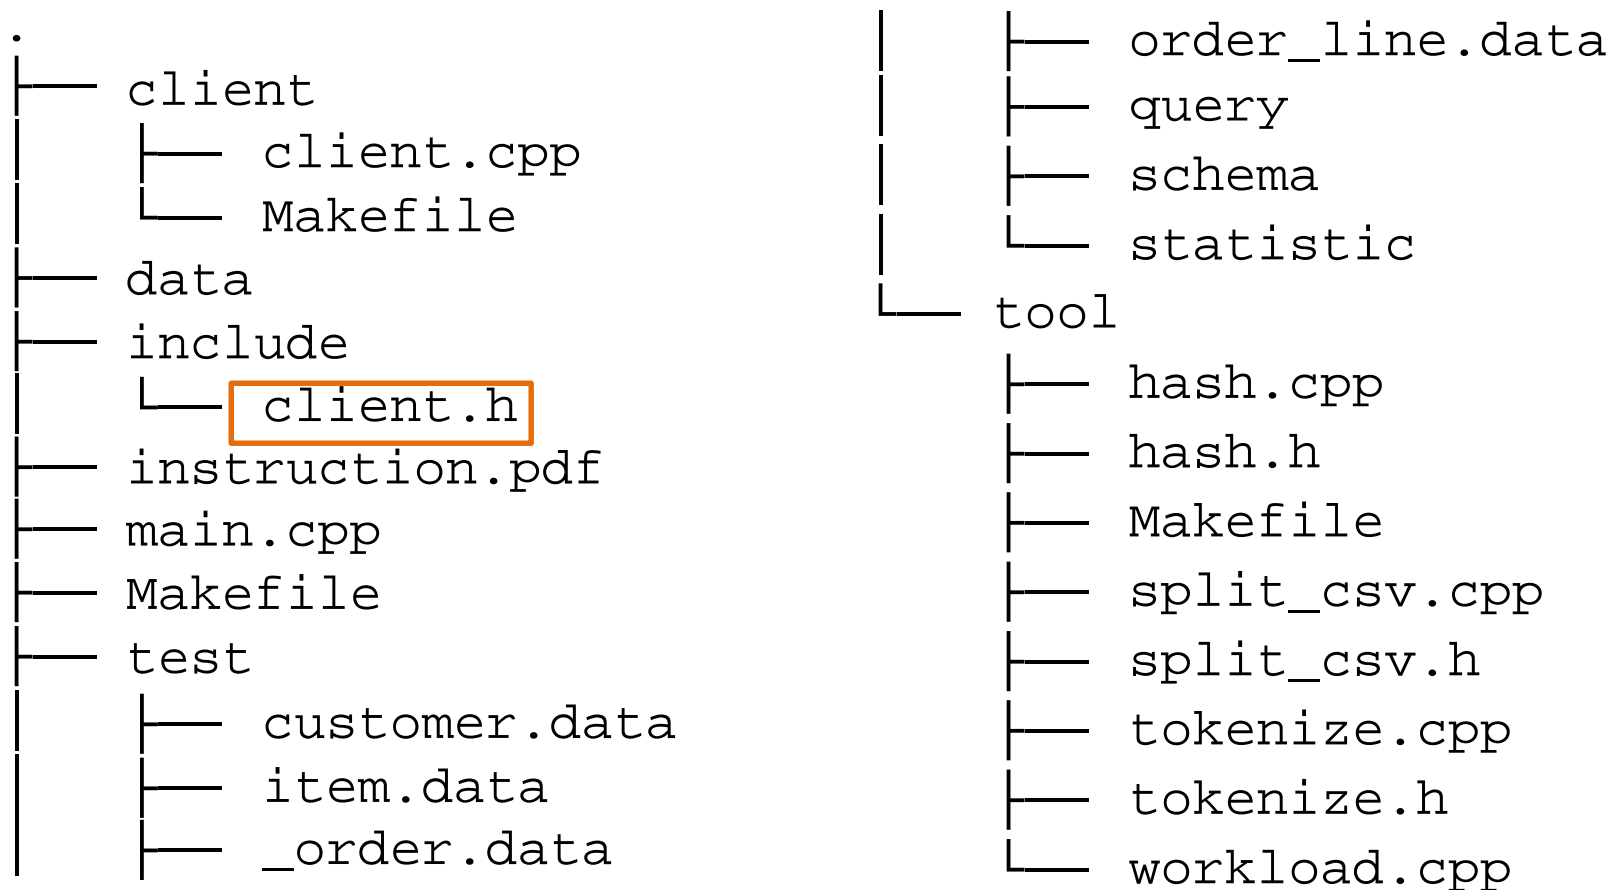
Nonnegotiable



# Directory Structure



# Directory Structure



# You Have to Implement

- **create()**

*Create a new table.*

- **train()**

*Given some query information, train your system and choose the storage and access methods.*

- **load()**

*Load initial data in csv format. The initial data set might be too large to keep in the main memory entirely.*

- **preprocess()**

*Build the indexes and do other preprocessing.*

<G:\助教\working\讲义&ppt\course\client\client.cpp>

# You Have to Implement

- `execute()`

*Execute a query or insert statement.*

- `next()`

*Get the next row from the result set of the last query.*

- `close()`

*Close the sockets and kill other threads.*

*See `course/include/client.h` for more details.*

# You Have to Implement

- **execute()**

*Execute a query or insert statement.*

- **next()**

*Get the next row from the result set of the last query.*


- **close()**

*Close the sockets and kill other threads.*

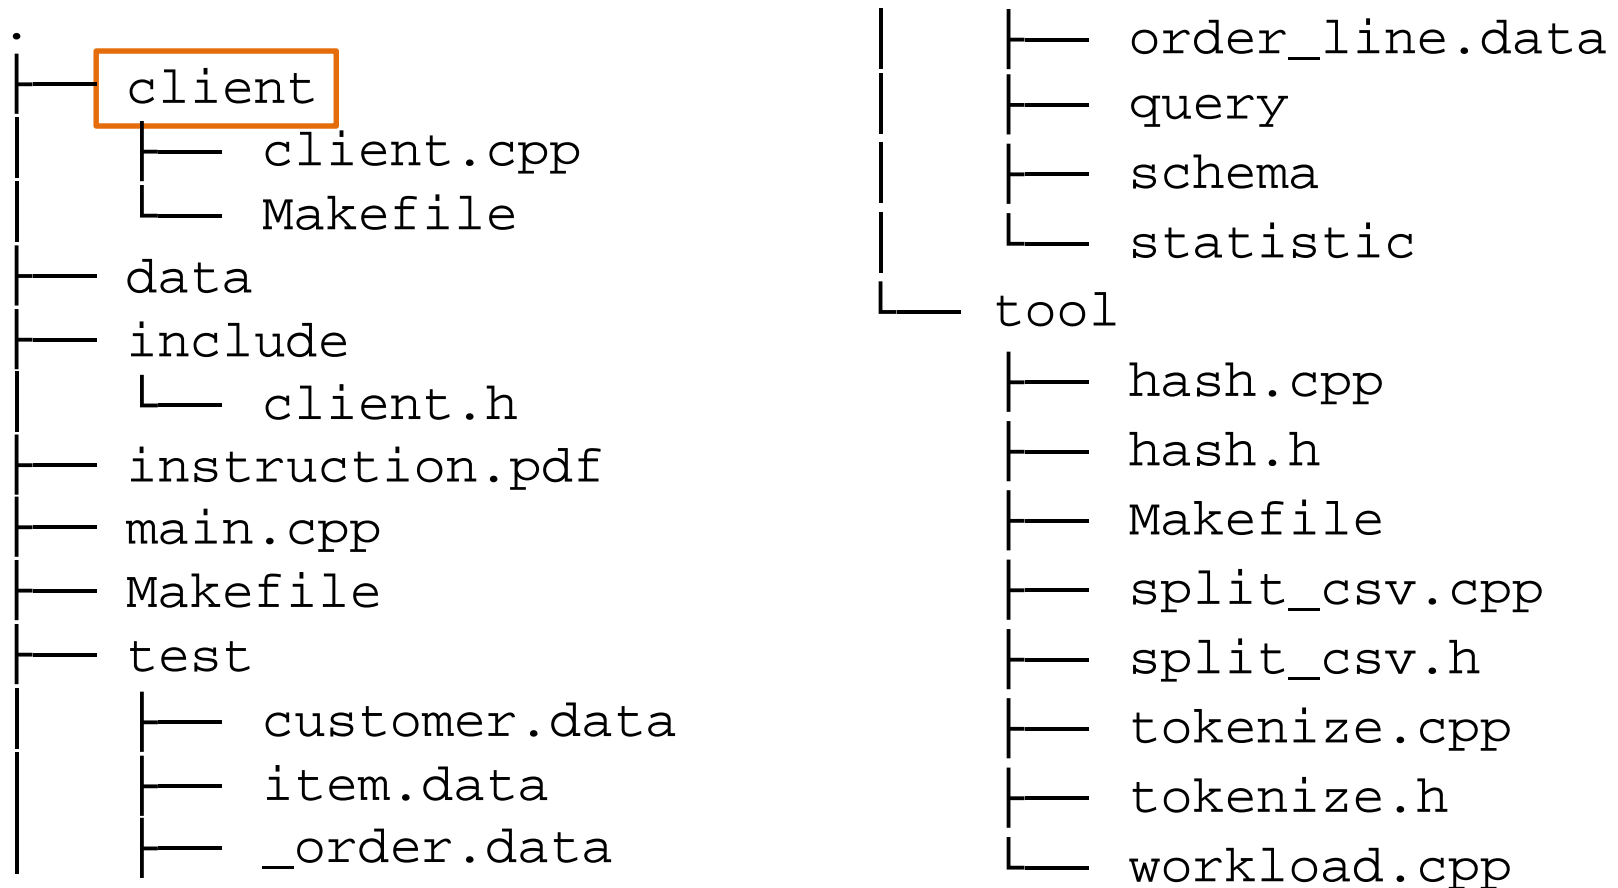
**WARNING: Run time of execute()  
and next() will be measured.**

*See course/include/client.h for more details.*

# Test Procedure

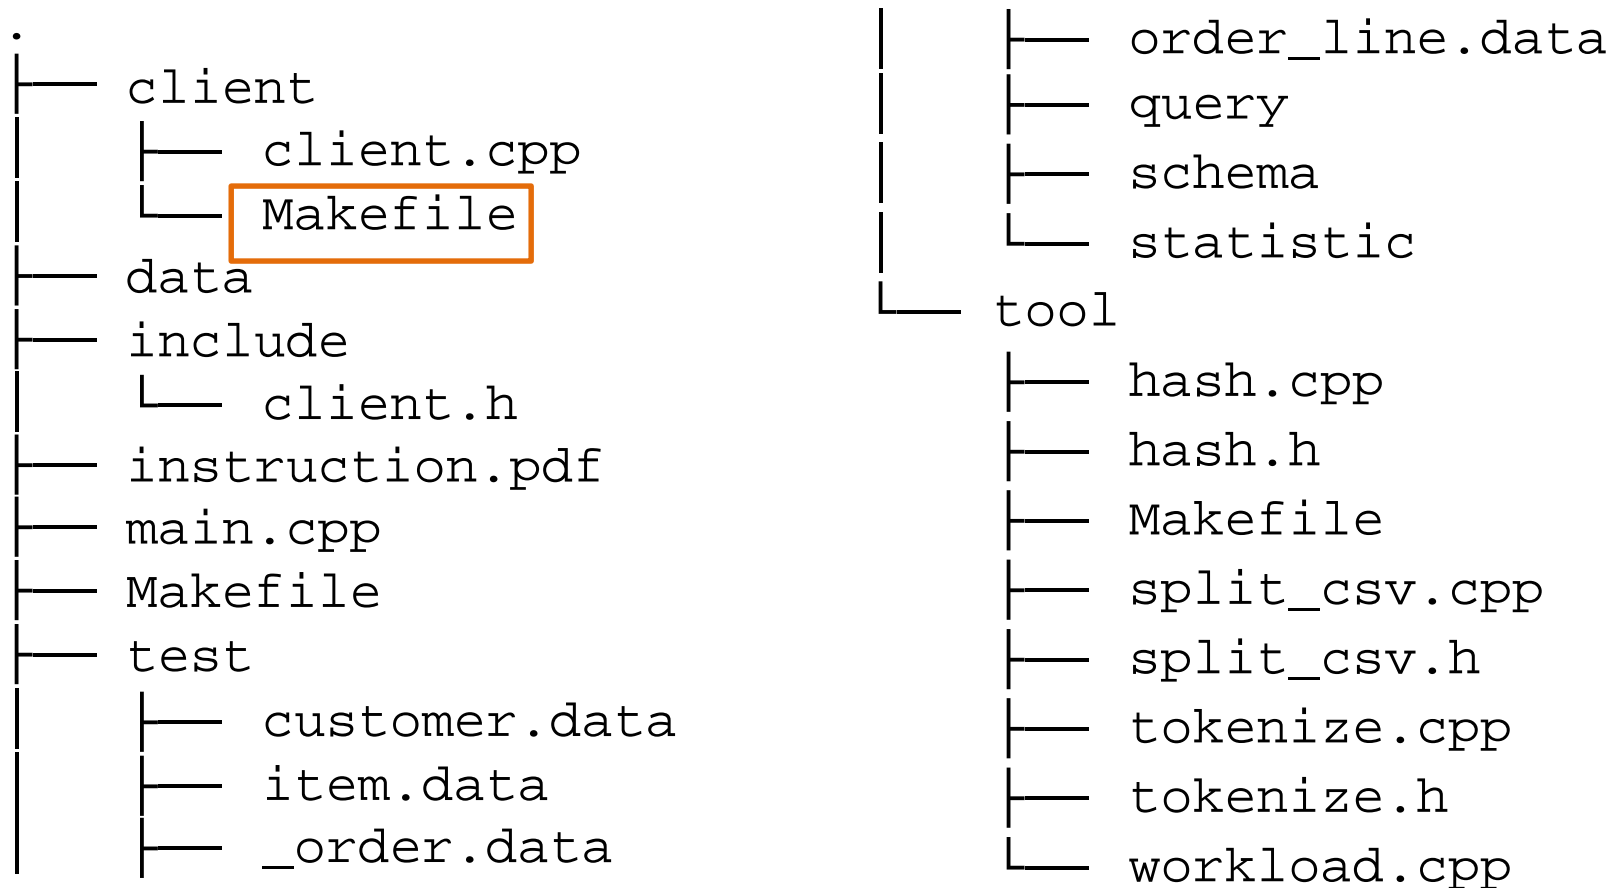


# Directory Structure



*Keep all your source codes in this directory.*

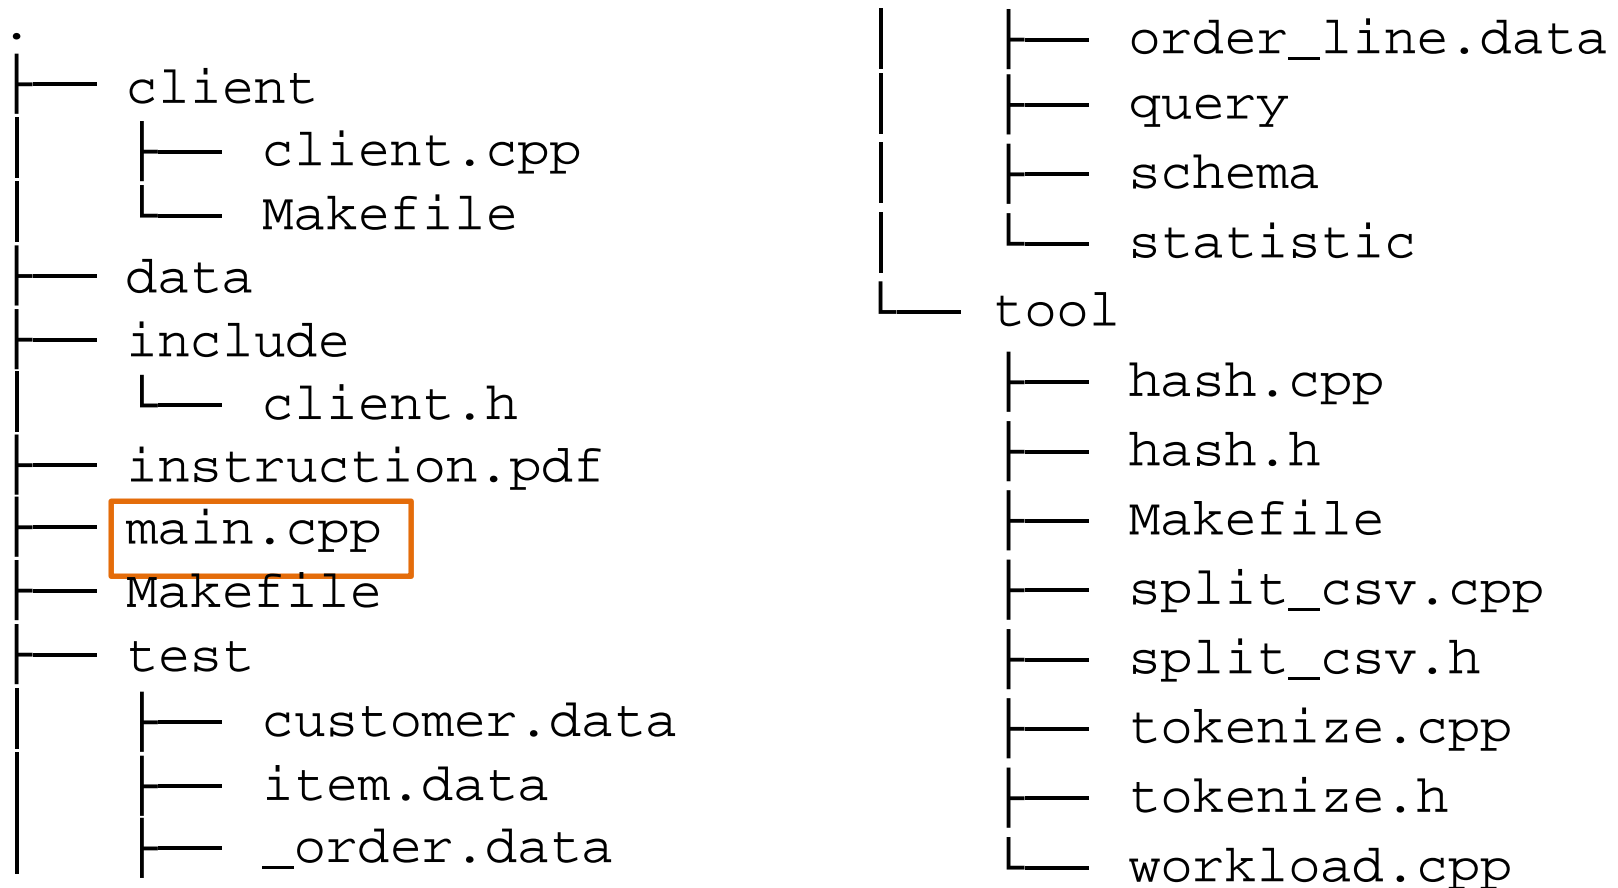
# Directory Structure



*Make sure your Makefile is correct.*

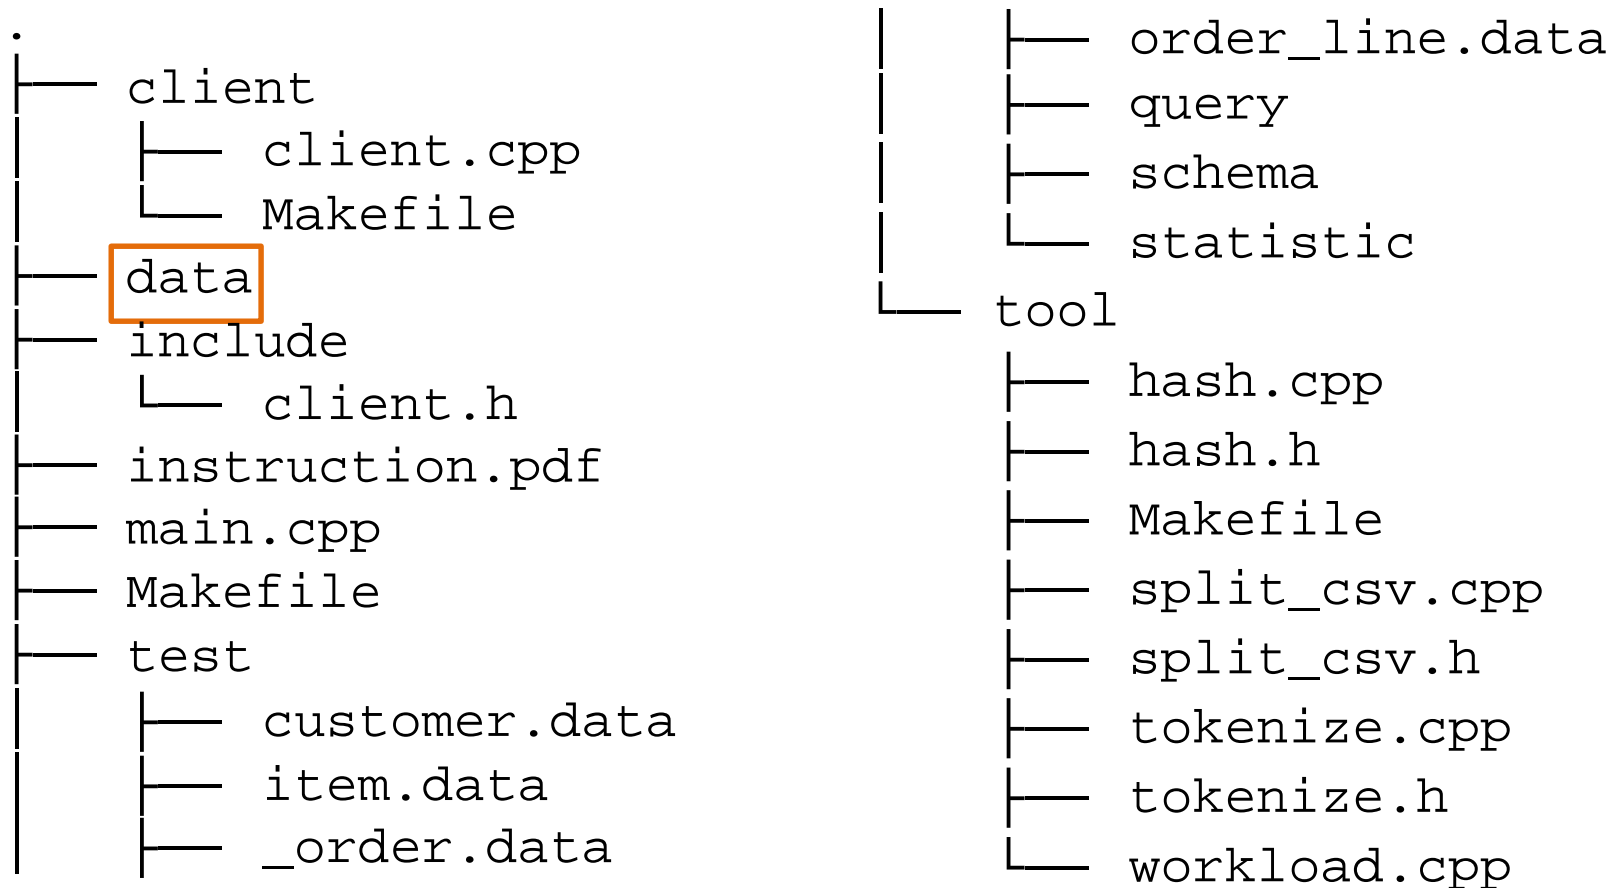


# Directory Structure



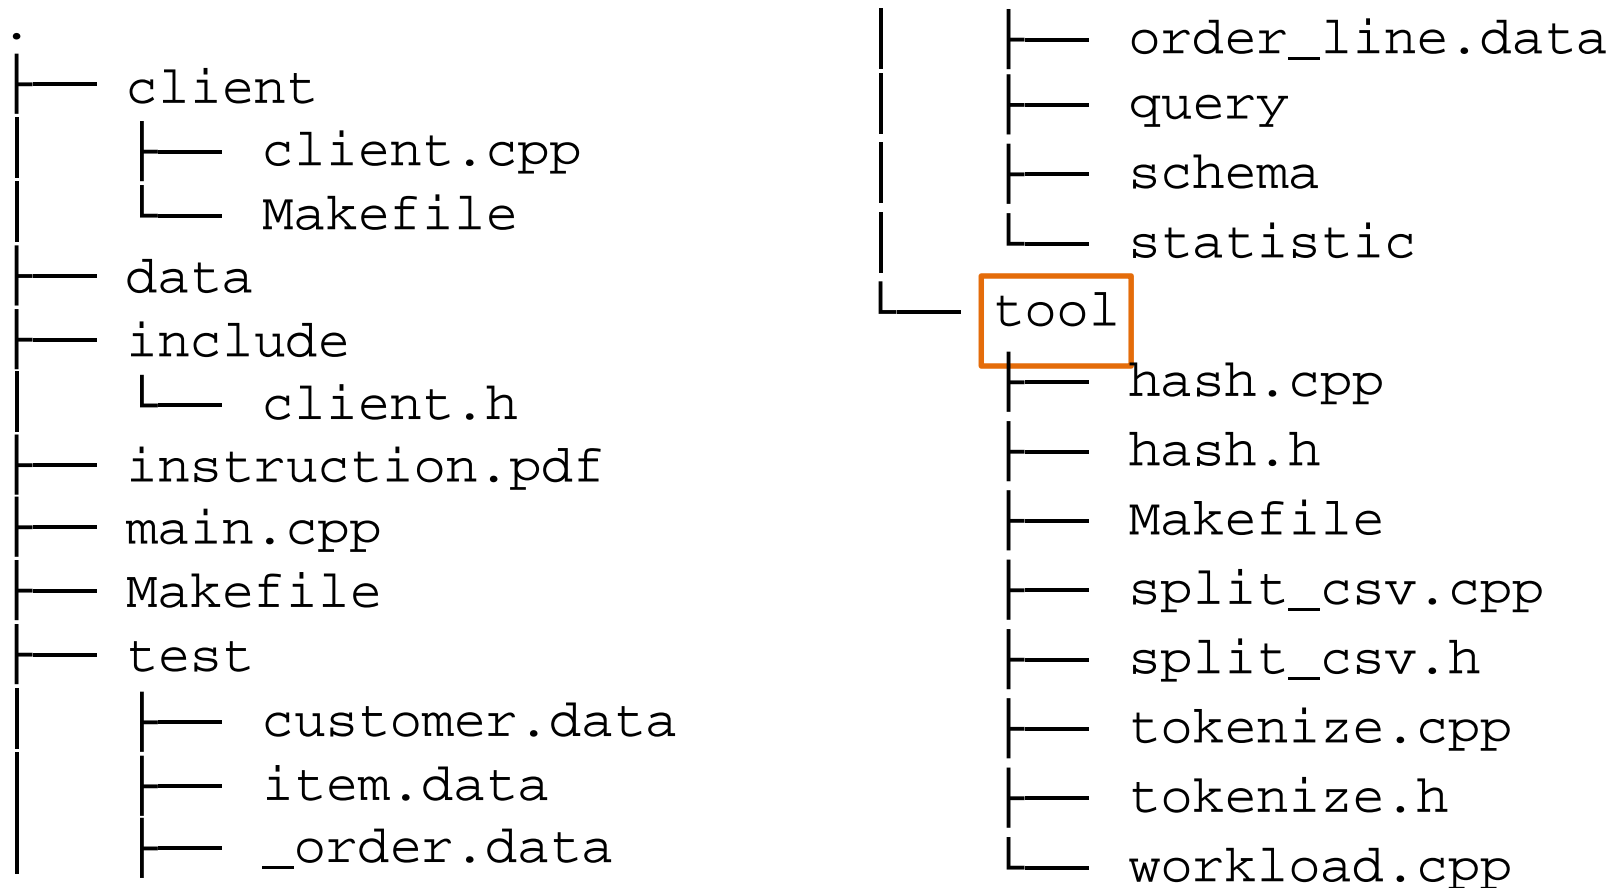
*Test procedure. This file will be modified.*

# Directory Structure



*Keep all your data in this directory.*

# Directory Structure



*Do NOT use any routines in other folders.*

# Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```

A condition could be

`column = constant`

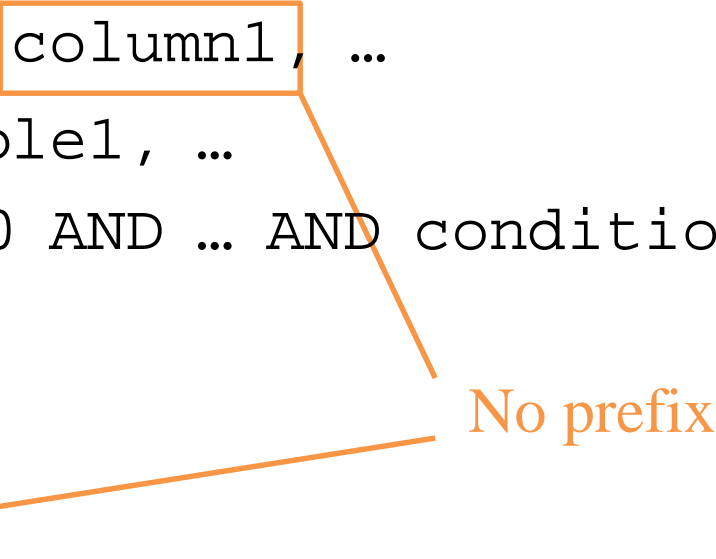
`column < constant` (For integers only)

`column > constant` (For integers only)

`column0 = column1` (Join condition)

# Query Statement

```
SELECT column0, column1, ...  
FROM table0, table1, ...  
WHERE condition0 AND ... AND conditionN;
```



A condition could be

No prefix

column = constant

column < constant (For integers only)

column > constant (For integers only)

column0 = column1 (Join condition)

# Query Statement

```
SELECT column0, column1, ...
```

```
FROM table0, table1, ...
```

```
WHERE condition0 AND ... AND conditionN;
```

A condition could be

The only operator

```
column = constant
```

```
column < constant (For integers only)
```

```
column > constant (For integers only)
```

```
column0 = column1 (Join condition)
```

# Query Statement

```
SELECT column0, column1, ...
```

```
FROM table0, table1, ...
```

```
WHERE condition0 AND ... AND conditionN;
```

A condition could be

```
column = constant
```

```
column < constant (For integers only)
```

```
column > constant (For integers only)
```

```
column0 = column1 (Join condition)
```

If the FROM-clause contains only one table, there might be no WHERE-clause.

# Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```

All value lists are in csv format.

```
constant0,constant1,...,constantN
```



# Insert Statement

```
INSERT INTO table  
VALUES (value_list0), ..., (value_listN);
```



All value lists are in csv format.

No column list

```
constant0,constant1,...,constantN
```

# Insert Statement

```
INSERT INTO table
```

```
VALUES (value_list0), ..., (value_listN);
```

All value lists are in csv format.

Number of rows is important  
for the `train()` routine.

```
constant0,constant1,...,constantN
```

# Format

```
SELECT _a_, _b_  
FROM _A_, _B_  
WHERE _a_ = _5_ AND _b_ < _10_;
```

```
INSERT INTO _A_ VALUES_  
(_0, 'Stalin', 1879_)_,_  
(_1, 'Roosevelt', 1882_)_;
```

# Data Types

- **INTEGER**

*32-bit unsigned integer, 'int' is OK.*

- **VARCHAR(d)**

*Consist of \_, a-z, A-Z, or 0-9. Enclosed by single quotes.  
At most d characters (excluding the quotes).*

## *NOTE:*

*All identifiers (table names and column names) are string constants not starting with 0-9.*

*Columns in different tables have distinct names.*

*String constants don't contain space, quote, or comma.*

# Primary Keys

- Primary keys will be assigned to all relations.
- The primary keys will be unique. There is no need to check this constraint.
- The primary keys will be given in ascending order.
- You can just ignore them.

# Join Operations

Let nodes represent tables and edges represent join conditions, then each query can be transformed into a graph. This graph should be a tree which

- is connected;
- contains no self-cycles;
- contains no duplicate edges;
- contains no cycles (at least 3 nodes).

# Workloads

- Projection
  - Selection
  - Join
- 
- TPC-C
  - TPC-H

# About Third-party Library

- You are free to use any third-party library or code about storage, index, multi-thread, network, etc.

*e.g. Boost, Berkeley DB, open-source disk-based B-tree / hash table implementation, etc*

- You are forbidden to use any system that is capable to process a SQL query.

*e.g. MySQL, PostgreSQL, etc*

- Ask for confirmation if you are not sure.



# Document / Presentation

- System Architecture
- Storage Model and the Selection Strategy
- Index Structure and the Selection Strategy
- Caching Strategy
- Query Processing Strategy
  - Heuristic Rules
  - Cost Model
- Other features of your system
- References
- Personal Contribution Rate (For document)

# Submission

- Prepare your submission with *make tar*.
- Submit the *\*.tar.gz* to ...
- One submission every 2 hours.
- Only the last submission counts.

# Notice

- Some tests may depend on the result of another one.
- Time limitation applies to the whole program, although only `execute()` and `next()` affect your final scores.
- Not all the queries are provided in `train()`, although only those included are measured.

# Warnings

- Never do irrelevant operations
- Never replicate other team's work



# Hints

- Read some books and research papers
- Discuss with others
- Start **ASAP**

create()

Keep the schema safe.

# train()

- Find affinitive tables.
- Find affinitive attributes.
- Choose access methods.
- Read-intensive or update-intensive?

load()

Keep the data safe.



# preprocess()

- Make some useful statistics.
- Build some indexes.
- Start some threads.

# Statistics

- For uniform distribution
  - $\text{Size}(R)$ ,  $\text{Cnt}(R)$ ,  $\text{Card}(A)$ ,  $\text{Min}(A)$ ,  $\text{Max}(A)$
  - $\text{SF}(A = \text{value}) = 1 / \text{Card}(A)$
  - $\text{SF}(A > \text{value}) = (\text{Max}(A) - \text{value}) / \text{Range}(A)$
  - $\text{SF}(A < \text{value}) = (\text{value} - \text{Min}(A)) / \text{Range}(A)$
  - $\text{SF}(A_0 \wedge A_1) = \text{SF}(A_0) * \text{SF}(A_1)$
- For skewed distribution (e.g., Zipf)
  - Histogram

# execute() and next()

- Do all the jobs in execute().
- Do all the jobs in next().
- Do all the jobs in independent threads.

# Query Processing

- Google 'query processing'
- Search on ACM Digital Library (*dl.acm.org*)
- Cost model vs. Rules


# Join Operation

- Nested Loop Join
- Index-based Nested Loop Join
- Sort-Merge Join
- Hash Join (Pruning)

close()

Close your program safely.







大学计算机教育国外著名教材、教参系列 (影印版)

# Principles of Distributed Database Systems

PEARSON  
Prentice  
Hall

Second Edition

M. Tamer Özsu  
Patrick Valduriez

## 分布式数据库 系统原理

(第2版)



清华大学出版社



# What to do immediately

- Homework: Chapter1-5
- Team partners: 2-3
  - Before 15<sup>th</sup> Oct.
  - Send me partner name+password
  - stoneliu2010@gmail.com

# Contact Information

- 办公室： 东主楼 10区204
- 刘思彤
  - 189 1163 5234
  - stoneliu2010@gmail.com
- 张晓航
  - 158 1150 7626
  - zhangxhscut@gmail.com

Good luck and have fun!